

Solving the Consensus Problem in a Mobile Environment

Nadjib BADACHE
USTHB
16111 Bab-Ezzouar, Algeria
badache@ist.cerist.dz

Michel HURFIN *
IRISA - INRIA
35042 Rennes, France
hurfin@irisa.fr

Raimundo MACÊDO†
LaSiD-CPD-UFBA
40170-110 Salvador BA, Brazil
macedo@ufba.br

Abstract

Atomic Broadcast or Non-Blocking Atomic Commitment are classic agreement problems encountered when designing fault-tolerant distributed systems. Specific protocols that solve such agreement problems can be designed based on a common building block, namely the consensus service. Unfortunately, the consensus problem has no deterministic solution in an asynchronous distributed system that is subject to even a single process crash failure. Among the solutions proposed to circumvent this impossibility result, the concept of unreliable failure detectors proposed by Chandra and Toueg is particularly attractive. They have defined a protocol that solves the consensus problem when the assumption that the underlying failure detector belongs to the class $\Diamond S$ holds true. This paper presents an extension of their protocol that allows to solve this fundamental problem in a mobile environment. In such an environment, the problem is more challenging: based on their initial states, a set of mobile hosts must agree on a common decision despite disconnections, changes of location and failures of mobile/fixed hosts.

1 Introduction

The wide use of portable computers and the advances in wireless networking technologies have greatly enhanced mobile computing which is now a major trend in both the computer and telecommunication industries. Intrinsically, a mobile host (*i.e.*, a portable, handheld or embedded computer) changes its location periodically. The connection to the network is usually temporary with periods of (voluntary/involuntary) disconnection. Mobile systems are often subject to environmental adversities which can cause loss of messages or data [7]. A mobile host can crash or suffer from frequent and intermittent disconnections from the rest of the network. Thus, designing fault-tolerant distributed applications in such environments is a complex

endeavor.

In recent years, several paradigms have been identified to simplify the design of fault-tolerant distributed applications in a conventional static system. The Consensus paradigm is one of the most fundamental since it abstracts other agreement problems. Given a fixed set of processes, the consensus problem is defined as follows: each process proposes an initial value to the others and, despite failures, all non-crashed processes have to agree on a common decision value, which depends on the initial proposals. Any solution to this basic problem can be used to solve other problems such as non-blocking atomic commitment or atomic broadcast [2, 6]. Non-blocking atomic commitment requires all participants in a transaction to take the same decision, namely COMMIT or ABORT the transaction. Atomic broadcast allows processes to agree on both a set of messages and a single delivery order for these messages. In both examples, a consensus service can be used as a basic building block. The semantics associated with the proposed and decided values differ from one agreement problem to another. Yet, the difficulty is always the same: the decision should not be postpone forever (termination property) and has to be unanimous (agreement property). Due to its wide applicability, the consensus problem has been extensively studied. Unfortunately, this problem has no deterministic solution in an asynchronous distributed system that is subject to even a single process crash failure [4]. Among the solutions proposed to circumvent this impossibility result, the concept of unreliable failure detectors proposed by Chandra and Toueg is particularly attractive [2]. In this approach, the weakest conditions that have to be satisfied to solve the consensus problem have been identified [3] and accordingly several protocols have been proposed to solve the consensus problem. These protocols are in no way trivial: this is due to the fact that they do not require reliable failure detectors.

Despite its usefulness, no work has been devoted to this problem in a mobile computing environment (to our knowledge). Unfortunately, existing protocols

*Supported by a France Télécom/CNET Grant 981B123

†Supported by CNPq Grants 300013/87-6 (Project BCG) and 680067/95-7 (Project LOCUS/ProTeM-CC)

are not suited to a mobile environment: the consensus problem is even more challenging in such an environment. The aim of this paper is to identify the inadequacies of existing protocols and to propose a solution to the consensus problem in a mobile computing environment. The paper is organized as follows: in Section 2, protocols that solve the consensus problem in a conventional asynchronous distributed system, are discussed. Section 3 describes the mobile system model and presents a protocol to solve the consensus problem in a mobile environment. We conclude in Section 4.

2 Consensus in a Static System

2.1 The Consensus Problem

We consider an asynchronous distributed system consisting of n processes denoted p_1, p_2, \dots, p_n . Processes communicate and synchronize by sending and receiving messages through channels. The distributed system is asynchronous: no assumptions are made as to the relative speed of processes or message transfer delays. Each pair of processes is connected by a reliable link. A process may fail solely by crashing, *i.e.*, by prematurely halting; it behaves correctly (*i.e.*, according to its specification) until it possibly crashes. By definition, a correct process is a process that does not crash during the course of an infinite run.

In the consensus problem, all correct processes have to reach a common decision on some value v , which must belong to the set of proposed values. The consensus problem is defined in terms of two primitives called *propose* and *decide*. Initially each process p_i selects a value v_i from a set of possible values and invokes the primitive *propose* with this value as a parameter: we say that p_i proposes v_i . A process ends its participation in the consensus by executing *decide*(v): we say that it decides the value v . Formally the following properties have to be held:

Termination: Every correct process eventually decides some value.

Agreement: No two processes decide differently.

Validity: If a process decides v , then v was proposed by some process.

In [4], Fischer, Lynch and Paterson have shown that consensus cannot be solved deterministically in an asynchronous system that is subject to even a single crash failure. To overcome this impossibility result, Chandra and Toueg propose to augment the asynchronous model of computation with the concept of *unreliable failure detectors* [2].

2.2 Unreliable Failure Detectors

A distributed failure detector is a set of n failure detector modules, one per process. The failure detector

module attached to p_i is an oracle in charge of giving hints about processes suspected to be faulty: it maintains a list *Suspected_i* containing the identities of processes it currently suspects to have crashed. The failure detector is referred to as unreliable since each module can make mistakes by erroneously adding or removing processes to its list of suspects.

A failure detector is defined in terms of two abstract properties, namely *completeness* and *accuracy*. Roughly speaking, completeness requires that a failure detector suspects every process that actually crashes, while accuracy restricts the mistakes that a failure detector can make. In [2], Chandra and Toueg define eight classes of failure detectors, depending on the nature of the completeness and accuracy properties. In this paper, we are interested in the $\Diamond S$ class of failure detectors. This class of failure detectors which has been proved to be sufficient to solve consensus [2], and even to be the weakest one [3], is specified by the following two properties. The *Strong Completeness* property states that eventually every process that crashes is permanently suspected by every correct process. The *Eventual Weak Accuracy* property states that there is a time after which some correct process is never suspected by any correct process.

2.3 Consensus Protocols based on $\Diamond S$

Several protocols designed to work with $\Diamond S$ failure detectors have been proposed [2, 5, 8]. They all require that a majority of processes is correct. They all are based on the rotating coordinator paradigm and proceed in consecutive asynchronous rounds. Each round r is coordinated by a predetermined process p_c defined by $c = (r \bmod n) + 1$. Thus, processes deal with a crash of the current coordinator by moving to the next round. The accuracy property of the failure detector ensures that there will be eventually a round during which the coordinator will not be suspected. The number of rounds performed by each process is arbitrary: it depends on the occurrence of failures and also on the behavior of the failure detector modules. Consequently, it is possible that not all the processes decide in the same round. So, in each protocol, a specific locking mechanism ensures there is a single decision value.

The first protocol which has been proposed (called CT in the following) is due to Chandra and Toueg [2]. While CT uses a centralized scheme (all messages are to/from the coordinator), the two other protocols use a decentralized scheme (each process sends messages to all processes). In all the protocols, the coordinator of round r tries to impose a particular value as the decision value. In CT, this value is not necessarily the estimate of the coordinator at the time it starts round r : on the contrary, this value is computed after the coordinator has gathered estimates from other processes.

For this reason, this protocol (unlike the two others) can be extended to solve a slightly different problem (See Section 2.4). This quality persuades us to select CT as the basis of the proposed solution.

In CT, each process p_i manages a local variable that represents its current estimate of the decision value (initially, the value of this estimate is equal to the initial value v_i proposed by p_i). During the execution of the successive rounds, this value is updated and converges to the decision value. More precisely, each round is divided into 4 phases:

- (1) In the first phase, each process sends to the current coordinator its own estimate of the final value.
- (2) The second phase is only executed by the coordinator. The coordinator gathers estimates from a majority of processes: it is assured to receive at least a majority of estimates, because a majority of processes is correct by assumption. Then it selects the estimate whose timestamp is the greatest one¹ and suggests this estimate by sending it to all the processes.
- (3) In the third phase each process p_i waits for the receipt of a new estimate from the coordinator. Either p_i suspects the coordinator to have crashed or p_i receives and adopts the new estimate. In the former case, a process sends a negative acknowledgment to the coordinator. In the latter case it sends a positive acknowledgment and updates the timestamp associated with its new estimate by setting it to the current value of its round counter.
- (4) The fourth phase is only performed by the coordinator. It waits for a majority of acknowledgment messages. If it receives only positive acknowledgments, it reliably broadcasts a decision message (otherwise, the coordinator proceeds to the next round). Informally, Reliable Broadcast guarantees that (a) all correct processes deliver the same set of messages, (b) all messages broadcast by correct processes are delivered, and (c) no spurious messages are ever delivered.

Note that an estimate is irremediably locked as soon as a majority of processes have sent a positive acknowledgment to the coordinator: then, no other value can be selected to be the final decision. When a process terminates round r , it immediately proceeds to round $r + 1$, except if it has received a decision message with the value v . In this case, the process decides the value v and terminates.

2.4 Towards an Extended Consensus Problem

As formulated, the consensus problem is a pure agreement problem. Each process proposes a credible outcome to the consensus service which then forces the adoption of one of these values. Before launching a consensus, a process has to compute the initial

value it will proposed, bearing in mind that this value will perhaps become the decision value. Thus, to solve practical agreement problems, processes must generally execute a preliminary exchange phase before executing the consensus protocol. During this phase each process p_i broadcasts relevant local information and then waits until it either receives information on any process p_k or it suspects p_k . At the end of this phase, each process has its own global view of the current global state and can compute the initial value it proposes to consensus. For example, consider the Non Blocking Atomic Commitment problem which assures that all correct participants in a transaction adopt the same decision, namely COMMIT or ABORT the transaction. During the exchange phase, each process broadcasts its local decision (a YES vote or a NO vote) to the other processes. Then, if it has received a YES vote from each process, it proposes COMMIT to the consensus. If it has received a NO vote or if it has not received a vote from some process that it suspects to have crashed, it proposes ABORT to the consensus.

As investigated in [6], the exchanged phase can be suppressed (and consequently the total number of exchanged messages can be reduced) by adopting another practical building block which extend the original consensus problem defined in Section 2.1. The difference between the extended consensus problem and the original one lies in the validity property. The decision is no longer a value proposed by a process but a collection of values proposed by different processes. In this paper, we adopt the following validity property: if a process decides a set of values V , then the set V contains only initial values proposed by processes and its cardinality is at least equal to α ($1 \leq \alpha \leq n$). The α parameter is set by the upper layer application program when it requires that a minimal number of processes participate in the underlying agreement protocol. Of course, any protocol used to solve this extended problem do not necessarily terminate if more than $n - \alpha$ processes have crashed definitively. As shown in [6], slight modifications of CT allow this new problem to be solved. The extended protocol does not require a preliminary exchange phase: the global view is computed by the coordinator while it gathers information from each processes during the second phase of a round. Unfortunately, solving the same problem in an mobile environment is more challenging. Before explaining the inadequacies of the above solution, we briefly describe the particularities of a mobile system.

3 Consensus in a Mobile System

3.1 The Mobile System Model

A mobile system is a distributed system consisting of two distinct sets of entities: a set of mobile hosts (MHs)

¹Initially every estimate is timestamped with 0.

and a set of fixed hosts referred to as Mobile Support Stations (MSSs). The MSSs and the communication paths connecting them form a static distributed system which is similar to the system described in Section 2.1.

A cell is defined as the geographical area covered by a MSS. A MSS serves as a base station if it is able to communicate with the MHs located within its cell via a wireless medium. A fixed host which is not a base station compares with a base station whose cell is never visited by mobile hosts. A MH can directly communicate with a MSS (and vice versa) if and only if this MH is located within the cell serviced by the MSS. In order to send messages to another MH that is not in the same cell, the source MH has to contact its local MSS which forwards the messages over the static network to the local MSS of the target MH. The receiving MSS, in its turn, forwards the messages over the wireless network to the target MH. When a MH moves from one cell to another, an *Handoff procedure* is executed by the MSSs of the two cells.

If its current base station fails by crashing, the connection between a MH and the rest of the system is broken. Yet the MH can reconnect to the network by moving into another cell covered by a correct base station. A MH may fail or voluntarily disconnect from the system. When a MH fails, its volatile state is lost.

The bandwidth of the wireless link connecting a MH to a MSS is significantly lower than bandwidth of the links between static hosts. In addition, MHs have tight constraints on power consumption relative to desktop machines, since they usually operate on stand-alone energy sources such as battery cells. Consequently, they often operate in a doze mode or voluntarily disconnect from the network. Transmission and reception of messages over wireless links also consume power at a MH. So, distributed protocols for mobile systems need to minimize communication over wireless links. Furthermore, MHs are less powerful than fixed hosts and have less memory and disk storage. Hence, while designing distributed protocols for mobile systems, the above constraints should be taken into account [1].

3.2 Why CT Needs to be Modified

In the following, we consider a broadcast group $G = (G_MSS, G_MH)$ of communicating mobile hosts, where G_MH is a set of m mobile hosts roaming in a geographical area (like a campus area) covered by a fixed set G_MSS of n base stations. The m mobile hosts are denoted h_1, h_2, \dots, h_m whereas the n base stations are denoted $MSS_1, MSS_2, \dots, MSS_n$. In so far, local mobile hosts of base station MSS_i will refer to mobile hosts that belong to G_MH and are currently locating in the MSS_i cell. In this environment, the consensus problem is defined over the set G_MH of mobile hosts. Each mobile host h_k proposes a value

v_k and the mobile hosts have to decide on a common value V which is a set of values proposed by at least α different mobile hosts ($1 \leq \alpha \leq m$). We assume that at least α mobile hosts will communicate their initial value. In other words, less than $m - \alpha$ mobile hosts have crashed definitively.

Due to the resources constraints of mobile hosts and the limited bandwidth of the wireless links, the proposed protocol has to be executed by the set of MSSs on behalf of the set G_MH of mobile hosts. We assume that the consensus is initiated by one or several mobile hosts which can be located in different MSSs. Without previously consulting the other mobile hosts, a mobile host requests that its current base station launches the consensus. The contacted base station reliably forwards the request to the other base stations. At the end of this initialization phase, either all (or none of) the correct base stations execute the rest of consensus protocol. Then the activity of a MSS is divided into three main subtasks: (1) a MSS interacts with mobile hosts located in its cell to collect their initial values. (2) a MSS interacts with other MSSs to agree on a subset of proposed values and (3) a MSS interacts with the mobile hosts located in its cell to communicate the final outcome. In our approach, a base station which participates in the consensus protocol, always acts on behalf of a subset of mobile hosts. More precisely, the value V_i proposed by a base station MSS_i is a collection of values proposed by mobile hosts. Initially, V_i contains only values from mobile hosts connected to MSS_i . After exchanging messages with other base stations, V_i will also include values from mobile hosts that have never moved into the cell of MSS_i . While the consensus is not completed, a base station builds up its collection until it contains values from at least α distinct mobile hosts. When a mobile host enters a new cell, the corresponding base station requests its initial value if the base station is not yet aware of it. The mobile host communicates this value even if it has already given this information to several other base stations: this new exchange is not useless because the base stations previously informed may have crashed.

As seen in Section 2.3, the locking mechanism used in CT relies on the assumption that a majority of processes is correct and participates in each round. From this point of view, mobility appears to be a major difficulty. Let us consider the following scenario: all the mobile hosts are located in the same cell. If the corresponding base station has not crashed, it collects α initial values and proposes this set of values to the coordinator during phase one of a round. The other base stations have to act on the behalf of no mobile hosts: they must participate in the consensus by proposing an estimate equal to the empty set (otherwise the protocol may block). More generally, a base station cannot

postpone the sending of its estimate to the coordinator until this estimate contains (possibly) α initial values. Yet, the protocol must allow a base station to communicate a more accurate estimate later.

Now assume that all the mobile hosts are in the cell of a crashed base station. They progressively move to cells managed by correct based stations. Unfortunately, these base stations can already have participated in a round r by sending an estimate equal to the empty set (See above). In that case, the coordinator of round r has possibly received a majority of estimates (all equal to the empty set) and already proposed the empty set as the new estimate. More generally, as the coordinator can be erroneously suspected by some base stations², it is not certain it will eventually obtained a new estimate (by piecing together the received proposals) containing at least α initial values. Yet, even if coordinators propose new estimates that are not acceptable, the protocol must ensure that the decision value contains at least α values.

This simple example gives an idea of how the original protocol is inadequate to cope with mobility. The next Section outlines the proposed solution.

3.3 The Proposed Protocol

The proposed solution is based on CT. The consensus is obtained after a sequence of asynchronous rounds. A round r is managed by the base station MSS_c such that $c = (r \bmod n) + 1$. Whereas CT assumes that a process always sends its estimate once per round (phase 1) and changes its estimate only when it adopts the value proposed by the coordinator (phase 3), our protocol partially removes these limitations. A base station is allowed to change its proposed value while this value does not reflect the decision of at least α mobile hosts. In other words, a base station can change its mind when it adds new values to its uncompleted collection of values. So, during phase 1 of a round, a base station may send up to $\alpha + 1$ messages to the coordinator. Furthermore, after receiving a new estimate from the coordinator, a base station can still send a negative acknowledgment, meaning that it does not agree with the proposed outcome result. More precisely, the value proposed by the coordinator is refused if it contains less than α mobile hosts values. As soon as a base station has gathered α values and sent a positive acknowledgment to a coordinator, its behavior is similar to that of a process as defined in CT. However, when a base station decides, it is nevertheless in charge of communicating the decision to the mobile hosts located in its cell.

²In that case, these base stations have proceed to the round $r + 1$ and will never communicate to the coordinator of round r more accurate estimates.

The protocol is structured into three parts. Figure 1 describes the role of an arbitrary mobile host h_k . Figure 2 presents the protocol executed by a base station MSS_i . Actions 4-9 are related to the interactions between a base station and its local mobile hosts (on one hand) and the rest of base stations (on the other hand). Actions 10-14 depicts the adapted Chandra-Toueg's protocol. Finally, the third part (see Figure 3) corresponds to the handoff protocol used to handle the change of location of the mobile hosts.

```
% Mobile host  $h_k$  is located in the cell of  $MSS_i$ .

(1) Upon the application requires to start a consensus
    send INIT.1 to  $MSS_i$ ;

(2) Upon receipt of INIT.3 from  $MSS_i$ ;
    % The initial value is provided by the application.
    send PROPOSE( $h_k$ , Initial.Value) to  $MSS_i$ ;

(3) Upon receipt of DECIDE(Decided.Value) from  $MSS_i$ ;
    % The decision value is delivered to the application.
```

Figure 1. Protocol Executed by a mobile

Local Context of a Mobile Host

- *Initial.Value* : Value provided by the application program running on a mobile host.

Local Context of a Base Station MSS_i

- *Local.MH_i* : Set containing the identities of the mobile hosts located in the cell of MSS_i .
- *Suspected_i* : Set containing the identities of the base stations suspected to be crashed. This list is managed by the local failure detector module of MSS_i .
- *r_i* : Sequence number which identifies the current round executed by MSS_i .
- *Phase_i* : Phase number in a round. When the protocol starts or ends, *Phase_i* is equal to 0. Otherwise this variable is either equal to 1, 2, 3 or 4.
- *State_i* : This variable is set to *decided* if the consensus has terminated. Otherwise it is set to *undecided*.
- *ts_i* : Sequence number of the last round during which a new estimate sent by a coordinator has been accepted as the new value of V_i .
- *P_i* : Set containing the identities of the mobile hosts whose initial values are already known by MSS_i . MSS_i collects values of the mobile hosts located in its cell until $|P_i| \geq \alpha$ holds (*End_collect_i = true*).
- *New.V_i* : Set containing the collected values.
- *V_i* : Last set of values proposed by MSS_i .
- *Log_i[r]* : Set containing the estimates received by the coordinator MSS_i during the r^{th} round.
- *Nb.P_i[r]* (resp. *Nb.N_i[r]*): Number of positive (resp. negative) acknowledgments received by the coordinator MSS_i during round r .

```

Maj := |G_MSS|/2; Phasei := 0; End_collecti := false;
New_Vi := φ; Vi := φ; Pi := φ;
Statei := undecided; ri := 0; tsi := 0;
∀ r ≥ 1 : Logi[r] := φ; Nb_Pi[r] := 0; Nb_Ni[r] := 0;
cobegin
(4) || Upon receipt of INIT_1 or INIT_2
    if (Phasei = 0) then
        send INIT_2 to all MSSs except MSSi;
        Phasei := 1;
        if (((Local_MHi ∩ G_MH) ≠ φ) ∧ (¬End_collecti)) then
            W_Broadcast INIT_3
        endif
    endif

(5) || Upon receipt of PROPOSE(hk, vk)
    if (¬End_collecti) then
        Pi := Pi ∪ {hk}; New_Vi := New_Vi ∪ {vk};
        if (|Pi| ≥ α) then End_collecti := true endif;
        if (Phasei > 1) then
            send ESTIMATE(MSSi, ri, New_Vi, Pi, tsi) to MSSc
        endif
    endif

(6) || Upon receipt of DECIDE(Vj)
    if (Statei = undecided) then
        Statei := decided; Vi := Vj;
        send DECIDE(Vj) to all MSSs except MSSi;
        W_Broadcast DECIDE(Vj);
        Phasei := 0
    endif

(7) || Upon receipt of ESTIMATE(MSSj, r, Vj, Pj, tsj)
    Logi[r] := Logi[r] ⊕ {(MSSj, r, Vj, tsj)};
    if (¬End_collecti) then
        Pi := Pi ∪ Pj; New_Vi := New_Vi ∪ Vj;
        if (|Pi| ≥ α) then End_collecti := true endif
    endif

(8) || Upon receipt of PA(MSSj, rj)
    Nb_Pi[rj] := Nb_Pi[rj] + 1

(9) || Upon receipt of NA(MSSj, rj)
    Nb_Ni[rj] := Nb_Ni[rj] + 1

(10) || Upon Phasei = 1
    ri := ri + 1; c := (ri mod n) + 1;
    if tsi = 0 then Vi := New_Vi endif;
    send ESTIMATE(MSSi, ri, Vi, Pi, tsi) to MSSc;
    if (i = c) then Phasei := 2 else Phasei := 3 endif

(11) || Upon (Phasei = 2) ∧ (|Logi[ri]| > Maj)
    Let tmax be the largest timestamp such that there exists
    at least one element (MSSj, rj, Vj, Pj, tsmax) ∈ Logi[ri]
    if (tmax > 0)
        then Select one of those particular elements
            denoted (MSSj, rj, Vj, Pj, tsmax);
            Vi := Vj
        else Vi := New_Vi
    endif;
    send NEW_EST(MSSi, ri, Vi, Pi, End_collecti) to all;
    Phasei := 3

(12) || Upon receipt of NEW_EST(MSSc, ri, Vc, Pc, End_collectc)
    if (Phasei = 3) then
        if (End_collectc)
            then Vi := Vc; tsi := ri; End_collecti := true;
            send PA(MSSi, ri) to MSSc
        else send NA(MSSi, ri) to MSSc;
            Pi := Pi ∪ Pc; New_Vi := New_Vi ∪ Vc;
            if (|Pi| ≥ α) then End_collecti := true endif
        endif;
        if (i = c) then Phasei := 4 else Phasei := 1 endif
    endif

(13) || Upon (Phasei = 3) ∧ (MSSc ∈ Suspectedi)
    send NA(MSSi, ri) to MSSc; Phasei := 1

(14) || Upon (Phasei = 4) ∧ (Nb_Pi[ri] + Nb_Ni[ri] > Maj)
    if (Nb_Pi[ri] > Maj)
        then send DECIDE(Vi) to all MSSs except MSSi;
        Statei := decided; Phasei := 0
    else Phasei := 1
    endif
coend

```

Figure 2. Protocol Executed by a Base Station

Messages

- **INIT_1:** Such a message is sent by a mobile host to its base station to initiate a consensus. See actions 1 and 4.
- **INIT_2:** When a base station is asked by a mobile host to initiate a consensus, it broadcasts this message to inform the other base stations that a consensus is started. To ensure a reliable broadcast of message INIT_2, each destination base station has to forward it to the other base stations. So, despite failures of base stations, all (or none) correct base stations will be aware that a consensus has been initiated. See action 4.
- **INIT_3:** This message is sent to a mobile host either when its base station is informed (on receipt of INIT_2 message) that a consensus has started or when the mobile host enters a new cell managed by a base station MSS_i which is not aware of its initial value and has not yet completed its collection of values ($|P_i| < \alpha$). See the handoff procedure and actions 4 and 2.
- **PROPOSE(-):** Such a message carries the value proposed by a mobile host to its local base station³. A base station MSS_i takes it into account if $|P_i| < \alpha$. See actions 2 and 5.
- **ESTIMATE(-):** This message carries the estimate proposed by a base station MSS_i to the current coordinator MSS_c. Each estimate is tagged with a timestamp ts_i identifying the round during which MSS_i has updated its estimate for the last time (see action 12). During round r, MSS_i sends a first ESTIMATE message during action 10. Other ESTIMATE messages can be sent during action 5 when MSS_i updates its collection of values. The estimates sent during round r to MSS_c (MSS_c is necessarily the coordinator of round r), are gathered and logged in a local buffer Log_c[r]. A base station MSS_i can propose multiple estimates during a round r but the coordinator MSS_c keeps only the most recent estimate sent by MSS_i during round r. The statement $Log_c[r] := Log_c[r] \oplus (MSS_i, r, V_i, ts_i)$ executed in action 7 is equivalent to two successive operations: (1) $Log_c[r] := Log_c[r] \cup \{(MSS_i, r, V_i, ts_i)\}$ and (2) if there exists $(MSS_i, r, V'_i, ts_i) \in Log_c[r]$ and $(MSS_i, r, V''_i, ts_i) \in Log_c[r]$ such that $card(V'_i) \leq card(V''_i)$ then (MSS_i, r, V'_i, ts_i) is removed from Log_c[r]. While the collect is still possible ($End_collect_c = false$), the coordinator updates the sets New_V_c and P_c each time it receives the local view of another base station. See actions 10, 5 and 7.
- **NEW_EST(-):** This message carries the estimate proposed by the coordinator to the base stations. When the coordinator of round r has gathered a majority of

³The reader can notice that the value proposed by a mobile host is not required to be always the same. This possibility is not discussed in this paper.

estimates, it selects one estimate from its local buffer $Log_c[r]$ and sends it as a new estimate to all base stations. The selected estimate is either the new estimate sent by a previous coordinator which failed to gather a majority of positive acknowledgments or the set of values New_V_c of the current coordinator. While a base station MSS_i is waiting for a new estimate V_c , it asks its failure detector module whether the current coordinator has crashed or not. If the NEW_EST message is received before the coordinator is suspected and if it carries at least α participant mobile hosts, the base station updates its set of values V_i to V_c and replies with a positive acknowledgment. Otherwise it replies with a negative acknowledgment and next updates its sets New_V_i and P_i . See actions 11 and 12.

- PA(-): Positive acknowledgment sent to the coordinator. If the coordinator gathers a majority of positive acknowledgments, the set V_c is locked and broadcasted as the decided set of values to all base stations. Otherwise the coordinator moves to phase 1 and initiates the next round. See actions 12, 8 and 14.

- NA(-): Negative acknowledgment sent to the coordinator. See actions 12, 13, 9 and 14.

- DECIDE(-): This message carries the decided value. A base station MSS_i receives a message DECIDE(V_j) when a coordinator is aware that a majority of base stations agree upon the set of values V_j . MSS_i adopts this value, changes its state to *decided*, forwards the decided set of values to local mobile hosts and terminates. To ensure that all correct processes decide, the message is also forwarded to the other base stations (reliable broadcast). See the handoff procedure and actions 14, 6 and 3.

- GUEST(-): Such a message is sent by a mobile host to inform the current base station when it enters a new cell. See the handoff procedure.

- BEGIN_HANDOFF(-): A BEGIN_HANDOFF message is sent by MSS_i to MSS_j when MSS_i learns that a mobile host has moved from MSS_j cell to its own cell. See the handoff procedure.

- **Role of h_k**
Upon entry in MSS_i cell
send GUEST(h_k, MSS_j) to MSS_i
- **Role of MSS_i**
Upon receipt of GUEST(h_k, MSS_j)
 $Local_MH_i := Local_MH_i \cup \{h_k\}$;
send BEGIN_HANDOFF(h_k, MSS_i) to MSS_j ;
if $((Phase_i \neq 0) \wedge (h_k \notin P_i) \wedge \neg(End_collect_i))$ then
 send INIT_3 to h_k endif;
if $((Phase_i = 0) \wedge (State_i = decided))$ then
 send DECIDE(V_i) to h_k endif
- **Role of MSS_j**
Upon receipt of BEGIN_HANDOFF(h_k, MSS_i)
 $Local_MH_j := Local_MH_j - \{h_k\}$

Figure 3. Handoff Procedure

4 Conclusion

We have recall the interest of an extended consensus problem and shown how to overcome the difficulties induced by mobility when solving this problem. In an environment with m mobile hosts and n base stations, the proposed protocol tolerates up to $f' = m - \alpha$ mobile hosts failures and $f < \lceil n + 1 \rceil / 2$ base stations failures. The communications over wireless links are limited to a few messages (in the best case, two messages: one to propose the initial value and other to get the decided value). The mobile host's CPU time is low since the actual consensus is run by the base stations. The protocol is scalable: it is independent of the overall number of mobile hosts and all needed data structures are managed by the base stations. The threshold parameter α can be interpreted as a measure of the quality of the decided value delivered by the protocol. The value of parameter α is defined depending on the expected amount of mobile hosts failures that can be tolerated by a given application.

References

- [1] Badrinath B.R., Acharya A. and Imielinski T. Impact of mobility on distributed computations, *Operating Systems Review*, **27**(2), pp. 15–20, April 1993.
- [2] Chandra T.D. and Toueg S. Unreliable failure detectors for reliable distributed systems. *Journal of the ACM*, **43**(2), pp. 225–267, March 1996.
- [3] Chandra T.D., Hadzilacos V. and Toueg S. The Weakest failure detector for solving consensus. *Journal of the ACM*, **43**(4), pp. 685–722, July 1996.
- [4] Fischer M.J., Lynch N. and Paterson M.S. Impossibility of distributed consensus with one faulty process. *Journal of the ACM*, **32**(2), pp. 374–382, April 1985.
- [5] Hurfin M. and Raynal M. Fast asynchronous consensus with a weak failure detector. Technical Report #1118, IRISA, Rennes, July 1997 (submitted to publication).
- [6] Hurfin M., Raynal M. and Tronel F. A practical building block for solving agreement problems. In *Proc. IEEE IPCCC*, pp. 25–31, Phoenix, Feb. 1998.
- [7] Pradhan D.K., Krishna P. and Vaidya N.H. Recoverable mobile environments: Design and trade-off analysis. In *Proc. 26th IEEE Int. Symp. on Fault-Tolerant Computing Systems*, Sendai, Japan, June 1996.
- [8] Schiper A. Early consensus in an asynchronous system with a weak failure detector. *Distributed Computing*, **10**(3), pp. 149–157, 1997.